

# Lab 6

12 February 2024 18:56



lab6

## Lab 6

### Using DynamoDB with EC2 (Python based)

We have set up an EC2 instance and learnt how to write and execute Python code on our instance. In this lab, we will learn to write Python code to connect to a **DynamoDB** instance on AWS. We will learn how to create tables, add data into these tables, and query these tables, etc. We will write our code in several python scripts residing on the EC2 instance.

In the overall setup, your TCP client (on your local machine) can accept data from your IoT devices and relay it to the TCP server on EC2. The server in turn can invoke methods in the dbmanager to store this data into DynamoDB. Similarly, data may be retrieved from DynamoDB.

The python codes used in this guide are available [here](#).

## 1. Downloading and installing boto3: the AWS SDK for Python

Boto3 is the AWS [SDK](#) for Python. It makes it easy to integrate your Python application, library, or script with AWS services including Amazon S3, Amazon EC2, Amazon DynamoDB, and more. In our case, we will use Boto3 specifically for DynamoDB.

The following two commands should do the job:

```
$ sudo apt-get update  
$ sudo apt-get -y install python3-boto3
```

Here are some [alternative methods](#).

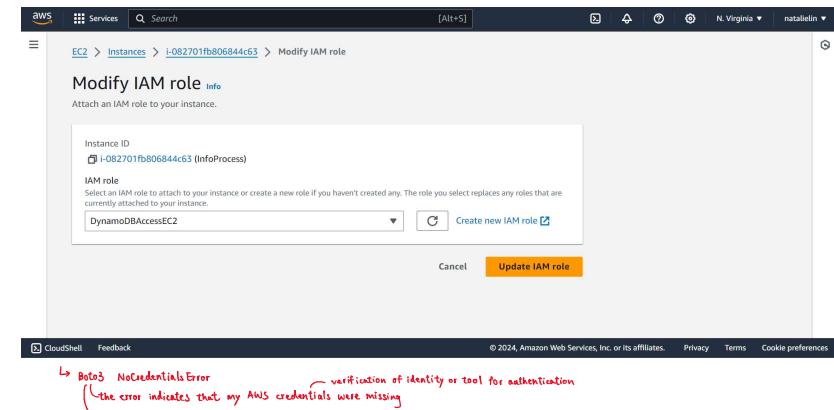
## 2. Creating a DynamoDB database

In this section, we'll carry out various operations on a DynamoDB instance. These include table creation, loading sample data, CRUD operations (create, read, update and delete an item), query and scan data, and delete table. The examples below have been taken from AWS's DynamoDB Developer Guide. Note that our focus here is to learn how to carry out these operations from within Python code. This is by no means the only way to access DynamoDB, and you may explore other ways if necessary, such as using JavaScript.

### 2.1 Create a table

Read the code below for creating a table called Movies. This is available in `MoviesCreateTable.py`

NosQL database  
less downtime - system should be able to survive the loss of one or more of these machines  
more scalable than SQL, horizontal scaling means greater overall capacity than vertical scaling (SQL)  
NoSQL is preferred for large and frequently changing data set  
increasing the capacity of a system by adding additional machines (nodes),  
as opposed to increasing the capacity of a existing machine (vertical scaling)  
key-value store  
key = string/integer, unique for the entire data  
value = can be anything  
distribution / partitioning with hash function - distribute the key-value pair based on the returned hash value



```

Creating a table in DynamoDB: an example in Python
import boto3 import botocore.client # the official AWS SDK for Python

def create_movie_table(dynamodb=None):
    dynamodb = boto3.resource('dynamodb', region_name='us-east-1') if 'dynamodb' is not provided, create a new DynamoDB resource using boto3
    if not dynamodb:
        dynamodb = boto3.resource('dynamodb', region_name='us-east-1') if 'dynamodb' is not provided, create a new DynamoDB resource using boto3
    region of our EC2 instance

    table = dynamodb.create_table(
        TableName='Movies',
        KeySchema=[# Primary key schema of the table
            {# the combination of the Partition Key and the Sort key
                'AttributeName': 'year',
                'KeyType': 'HASH' # Partition key
            },
            {# DynamoDB uses the partition key as the input to hash function. Result of hash
                # function determines the physical partition in which the item will be stored.
                'AttributeName': 'title',
                'KeyType': 'RANGE' # Sort key
            }
        ],
        AttributeDefinitions=[# define the attributes used in the key
            {# schema, specifying their types
                'AttributeName': 'year',
                'AttributeType': 'N' numeric
            },
            {
                'AttributeName': 'title',
                'AttributeType': 'S' string
            }
        ],
        ProvisionedThroughput={# Set the maximum amount of capacity, in terms of read/ writes
            'ReadCapacityUnits': 10,
            'WriteCapacityUnits': 10
        }
    )
    return table

if __name__ == '__main__': # checks if the script is being ran as the main program
    movie_table = create_movie_table() # calls the function above to create a DynamoDB table and assigns the resulting Table to the variable 'movie_table'
    print("Table status:", movie_table.table_status) # indicates whether the created table is active or not

```

The table is created in the usual way, as described in class. The line `dynamodb = boto3.resource('dynamodb', region_name='us-east-1')` provides an object-oriented interface to the DynamoDB service. As you can imagine, we can use the resource method to get interfaces to other AWS services as well. Here is more on `boto3 resources` and how to access them.

The `region_name` parameter in `resource` points to the region containing your EC2 instance and impacts the locations of partitions that will be assigned to DynamoDB items. You can read more on [AWS regions and zones](#) here.

In `create_table` we have a composite primary key, specified by both a Partition Key and a Sort Key. This means that we ~~can insert items which have the same partition key but different sort keys~~, and these items will be sorted on the same partition in ascending order of the Sort Key.

Primary key schema of the table.  
the combination of the Partition Key and the Sort key  
forms a unique primary key, which maps to a single value  
in the database.

Partition Key describes the partition in which the item  
will be stored.

Sort Key determines the order in which items will be stored  
in disk.

define the attributes used in the key  
schema, specifying their types

Set the maximum amount of capacity, in terms of read/ writes  
per second, that an application can consume from a Table

Letters 'N' and 'S' stand for numeric and string types respectively. For more [data types](#) see here.

Provisioned throughput is the upper-bound on the number of per-second reads and writes your application is allowed to make on this table. Any more read/writes per-second will be throttled. This helps DynamoDB price its services when used outside the free tier. You can read more about [provisioned throughput](#) here.

Now we can run the script to create the table:

```
$ python3 MoviesCreateTable.py
```

In order to see all existing tables, you can run the following python statement. It is perhaps more convenient to run this statement on the python command prompt rather than adding it to the script.

```
$ python3
>>> import boto3
>>> db = boto3.resource('dynamodb', region_name='us-east-1')
>>> print(list(db.tables.all()))
>>>exit()
```

#### NOTE

Try running the script one more time. You'll get an error stating that the table 'Movies' already exists. If you have multiple python functions, and multiple python scripts running on your EC2 instance, each accessing DynamoDB with resource('dynamodb'), they are all accessing the same instance of the database. In fact, your DynamoDB instance is also shared across your EC2 instances. If you created another EC2 instance, assigned it the IAM role with access to DynamoDB, and tried to create a 'Movies' table you would get the same error again.

## 2.2 Load sample data

Often it would be desirable to load a bunch of items into an existing table at once. This can be done by providing the items to boto3 in JSON format.

In this example, the data consists of a few thousand movies from IMDB. Following is the JSON format used:

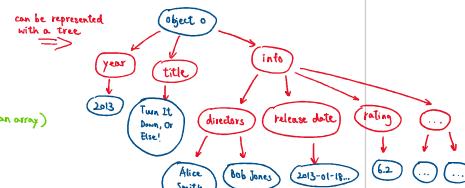
```
[  
  {  
    "year" : ... ,  
    "title" : ... ,  
    "info" : { ... }  
  },  
  {  
    "year" : ...,  
    "title" : ...,  
    "info" : { ... }  
  },
```

```
11      return self._make_api_call(operation_name, kwargs)  
File "/usr/lib/python3/dist-packages/botocore/client.py", line 635, in _make_a  
pi_call  
    raise error_class(parsed_response, operation_name)  
botocore.errorfactory.ResourceInUseException: An error occurred (ResourceInUseEx  
ception) when calling the CreateTable operation: Table already exists: Movies  
ubuntu@ip-172-31-21-227:~$ █
```

11

The year and title part are the primary key we specified while defining the table Movies. info consists of further information about the movies in JSON formation. Following is an example of what a single data item might look like:

```
object, denoted using curly braces { }  
key-value pairs are separated using commas ,  
string key stored with quotation marks  
"year" : 2013, colon separates key from its value  
"title" : "Turn It Down, Or Else!",  
"info" : {  
    "directors" : [ ordered array of values (objects can also be stored in an array)  
        "Alice Smith",  
        "Bob Jones"  
    ],  
    "release_date" : "2013-01-18T00:00:00Z",  
    "rating" : 6.2,  
    "genres" : [  
        "Comedy",  
        "Drama"  
    ],  
    "image_url" : "http://ia.media-imdb.com/images/I/09ERWAU7FS797A7L0UH09AMUP908RLlo5JF90EWR7LJKQ7@._V1_SX400  
"plot" : "A rock band plays their music at high volumes, annoying the  
"rank" : 11,  
"running_time_secs" : 5215,  
"actors" : [  
        "David Matthewman",  
        "Ann Thomas",  
        "Jonathan G. Neff"  
    ]  
}
```



Read more on the [JSON format](#) here.

semi-structured data  
each element can contain different amount of data - does not require predefined schema (unlike relational data model)  
text representation: good for exchange data between different apps, bad for performance (text takes up more space)

1. Download the sample data archive: [moviedata.zip](#)
  2. Extract the data file (moviedata.json) from the archive.
  3. Move the moviedata.json file to your EC2 instance using FileZilla.

Read the code below. This is available in `MoviesLoadData.py`

```
from decimal import Decimal imports the 'Decimal' class from the 'decimal' module so decimal numbers can be handled accurately
import json
import boto3

def load_movies(movies, dynamodb=None): function takes a list of movies and an optional DynamoDB resource
```

DynamoDB requires precise numeric representations, particularly when dealing with attributes of the data type 'Numeric'  
for indexing and sorting

```

if not dynamodb:
    dynamodb = boto3.resource('dynamodb', region_name='us-east-1')

table = dynamodb.Table('Movies') gets a reference to the 'Movies' table in Dynamodb

for movie in movies: iterates over each movie in the 'movies' list
    year = int(movie['year']) extracts the 'year' attribute from each movie and converts it to an integer
    title = movie['title'] extracts the 'title' attribute from each movie
    print("Adding movie:", year, title)

    table.put_item(Item=movie) function call to put an item(a movie) into the DynamoDB table

if __name__ == '__main__':
    with open("moviedata.json") as json_file: opens a JSON file named "moviedata.json" in the current directory
        movie_list = json.load(json_file, parse_float=Decimal) loads the JSONs data from the file into a Python list
ensures that floating-point numbers in the JSON file are parsed as 'Decimal' objects
    load_movies(movie_list) calls the function above to load data into the Table using
the dynamodb resource created (prevents loss of precision  
(e.g. rounding errors in Python)  
consistency in the data type)

```

Execute the script on the EC2 instance:

```
$ python3 MoviesLoadData.py
```

*NOTE: JSON is a good format in which to transmit data from a local computer to the server. The server can create a new record from the received JSON and insert it into a table. See below for how new items are created.*

## 2.3 Crud operations

Now the basic examples of creating, reading, updating and deleting items.

### 2.3.1 Create a new item

You can *add a new data item into an existing table by providing the primary key and the associated data*.

Read the code below. This is available in `MoviesItemOps01.py`

```

from pprint import pprint used for pretty-printing data structures, making output more readable
import boto3

def put_movie(title, year, plot, rating, dynamodb=None):
    if not dynamodb:
        dynamodb = boto3.resource('dynamodb', region_name='us-east-1')

    table = dynamodb.Table('Movies') refers to the 'Movies' table
    response = table.put_item(
        Item={
```

```

ubuntu@ip-172-31-21-227: ~
Adding movie: 2006 Breaking and Entering
Adding movie: 2013 Barbershop: The Next Cut
Adding movie: 2013 Knights of Badassdom
Adding movie: 1920 Das Cabinet des Dr. Caligari
Adding movie: 2012 Biohazard: Damnation
Adding movie: 1997 Fools Rush In
Adding movie: 2013 Bombshell
Adding movie: 2001 McCormick Jack
Adding movie: 2006 Les anges exterminateurs
Adding movie: 2013 Monica 2
Adding movie: 2006 The Sentinel
Adding movie: 2012 The Canarian Violet
Adding movie: 2012 Les infideles
Adding movie: 1995 Dolores Claiborne
Adding movie: 2011 The Entitled
Adding movie: 1990 The Legend of Navarone
Adding movie: 1989 Kickboxer
Adding movie: 2011 The Pill
Adding movie: 1966 Grand Prix
Adding movie: 2009 Sour Kitchen
Adding movie: 2012 Heartbreak Souls
Adding movie: 2012 The Attack
Adding movie: 2012 The Bates Haunting
Adding movie: 1956 Moby Dick
Adding movie: 2008 The Philadelphia Experiment
Adding movie: 1991 Dead Man on Campus
Adding movie: 2002 Slackers
Adding movie: 2008 From Within
Adding movie: 2010 Les aventures extraordinaires d'Adele Blanc-Sec
Adding movie: 1993 The Game Plan for Bobby Fischer
Adding movie: 2012 Sassy Pants
Adding movie: 1982 The Verdict
Adding movie: 1999 For Love of the Game
Adding movie: 2008 The Last Kiss With Me
Adding movie: 2008 The Accidental Husband
Adding movie: 1962 Cape Fear
Adding movie: 1957 Trooper Hook
Adding movie: 1949 Lord of the Flies
Adding movie: 1992 Cinema para Chocolate
Adding movie: 2012 Renoir
Adding movie: 2006 The Contract
Adding movie: 2010 The Clinic
Adding movie: 2004 Little Black Book
ubuntu@ip-172-31-21-227: ~
```

```

        'year': year,
        'title': title,
        'info': {
            'plot': plot,
            'rating': rating
        }
    }

    return response
}

if __name__ == '__main__':
    movie_resp = put_movie("The Big New Movie", 2015,
                           "Nothing happens at all.", 0)
    print("Put movie succeeded:")
    pprint(movie_resp, sort_dicts=False)

```

① nested attribute

② returns the response from the 'put\_item' operation, which contains information about the success of the operation

③ instructs 'pprint' not to sort the response from the 'put\_item' operation by keys - maintain original order

Run the script to create the item:

\$ python3 MoviesItemOps01.py

### 2.3.2 Reading an item

The previous script added the following item into the Movies table.

```
{
    year: 2015,
    title: "The Big New Movie",
    info: {
        plot: "Nothing happens at all.",
        rating: 0
    }
}
```

An item can be **read** from its table using its primary key value. The method used to accomplish this task is `get_item`. Read the code below to see how to do this. This code is available in `MoviesItemOps02.py`.

```

from pprint import pprint
import boto3
from botocore.exceptions import ClientError

def get_movie(title, year, dynamodb=None):
    if not dynamodb:
        dynamodb = boto3.resource('dynamodb', region_name='us-east-1')

    table = dynamodb.Table('Movies')

    try:
        response = table.get_item(Key={'year': year, 'title': title})
    except ClientError as e:
        print(e.response['Error']['Message'])
    else:
        return response['Item']

```

tries to get an item from the Dynamodb table using the 'get-item' method. It provides the 'key' parameter, which is a dictionary with the primary key attributes

④ if an error occurs during the 'get-item' operation, it catches the 'ClientError' exception, prints the error message, and returns 'None'

⑤ if operation is successful, returns the retrieved item, which is stored in 'response[Item]'.

```

ubuntu@ip-172-31-21-227:~ Adding movie: 2004 Little Black Book
ubuntu@ip-172-31-21-227:~$ python3 MoviesItemOps01.py
Put movie succeeded:
{'ResponseMetadata': {'HTTPHeaders': {'connection': 'keep-alive',
                                         'content-length': '2',
                                         'content-type': 'application/x-amz-json-1.0',
                                         'date': 'Tue, 13 Feb 2024 00:53:51 GMT',
                                         'server': 'Server',
                                         'x-amz-crc32': '2745614147',
                                         'x-amzn-requestid': '5FJE4D3J05TMU464QMH169Q2J7VV4KQNSO5AEMVJF66Q9ASUAAJG'},
                           'HTTPStatusCode': 200,
                           'RequestId': '5FJE4D3J05TMU464QMH169Q2J7VV4KQNSO5AEMVJF66Q9ASUAAJG',
                           'RetryAttempts': 0}}
ubuntu@ip-172-31-21-227:~$ python3 MoviesItemOps02.py
python3: can't open file 'MoviesItemOps02.py': [Errno 2] No such file or directory
ubuntu@ip-172-31-21-227:~$ python3 MoviesItemOps02.py
Get movie succeeded:
{'info': {'plot': 'Nothing happens at all.', 'rating': Decimal('0')},
 'title': 'The Big New Movie',
 'year': Decimal('2015')}
ubuntu@ip-172-31-21-227:~$ 

```

information about the success of the 'put\_item' operation

```
if __name__ == '__main__':
    movie = get_movie("The Big New Movie", 2015,)
    if movie: if the function returns a non-empty result (item successfully retrieved)
        print("Get movie succeeded:")
        pprint(movie, sort_dicts=False)
```

Note the use of exception handling. If the item does not exist, a proper error message is displayed. This helps in debugging the code.

Run the script:

```
$ python3 MoviesItemOps02.py
```

### 2.3.3 Update an item

An existing method can be updated using `update_item`. Once again, the `primary key` needs to be supplied. With `update_item` you can: `modify the values of existing attributes in item, add new attributes to the item, or remove attributes from the item.`

The following table shows the existing item and its required updated version.

```
#existing item
{
    year: 2015,
    title: "The Big New Movie",
    info: {
        plot: "Nothing happens at all.",
        rating: 0
    }
}
#required updated version
{
    year: 2015,
    title: "The Big New Movie",
    info: {
        plot: "Everything happens all at once.",
        rating: 5.5,
        actors: ["Larry", "Moe", "Curly"]
    }
}
```

We have `modified the values of plot and rating`. We have `added a new list attribute, actors`.

Read the code below which accomplishes this task. This code is available in `MoviesItemOps03.py`

```
from decimal import Decimal
from pprint import pprint
import boto3

def update_movie(title, year, rating, plot, actors, dynamodb=None):
    if not dynamodb:
        dynamodb = boto3.resource('dynamodb', region_name='us-east-1')
```

```

table = dynamodb.Table('Movies')

response = table.update_item(
    Key={
        'year': year,
        'title': title
    },
    UpdateExpression="set info.rating=:r, info.plot=:p, info.actors=:a", defines the update operation to be performed on the item
    ExpressionAttributeValues={
        ':r': Decimal(rating),
        ':p': plot,
        ':a': actors
    },
    ReturnValues="UPDATED_NEW" requests the updated attributes of the item to be returned and specifies that, in case of a successful update, the response should include the 'UPDATED_NEW' values.
)
return response

if __name__ == '__main__': main is executed only if the script is run directly
    update_response = update_movie(
        "The Big New Movie", 2015, 5.5, "Everything happens all at once.",
        ["Larry", "Moe", "Curly"])
    print("Update movie succeeded:")
    pprint(update_response)

```

The important part of the code is the parameter `UpdateExpression`. The content assigned to `UpdateExpression` describes the changes we wish to see in the specified item. As you can see, there is a specific format to the way we describe these changes. In short, we first write the expression, which includes some place holders (expression attributes), such as `:r`, `:p` and `:a` in this case. In the `ExpressionAttributeValues` parameter, we specify the values of these place holders. Read more here about this format and the [UpdateExpression](#) parameter.

In addition to `UpdateExpression`, a [ConditionExpression](#) may be used to specify a condition that must be satisfied for the update to take place.

Run the script:

```
$ python3 MoviesItemOps03.py
```

### 2.3.4 Delete an Item

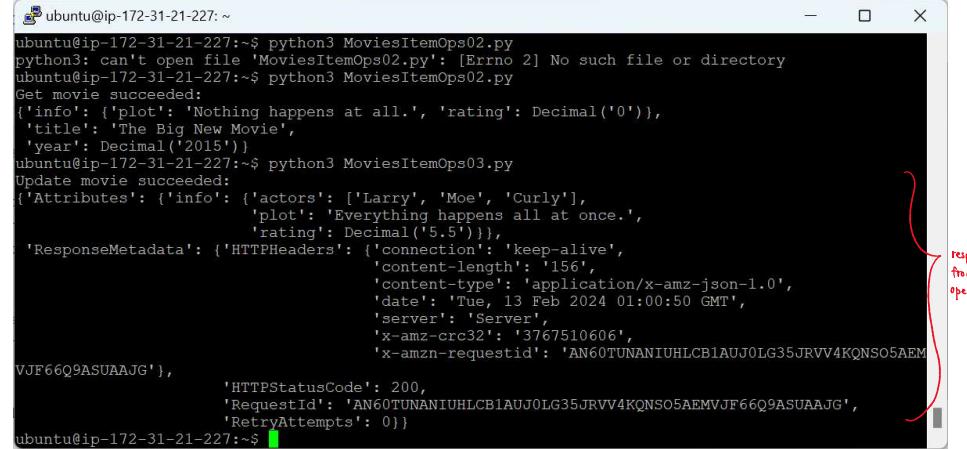
We use `delete_item` to delete an item using its [primary key](#). In the following example, you try to delete a specific movie item if its rating is `<=0` or less. Notice the use of `ConditionExpression` in this case. We don't always have to use `ConditionExpression`, but it's useful in many cases.

```

from decimal import Decimal
from pprint import pprint
import boto3
from botocore.exceptions import ClientError

def delete_underrated_movie(title, year, rating, dynamodb=None):
    if not dynamodb:

```



```

ubuntu@ip-172-31-21-227:~$ python3 MoviesItemOps03.py
python3: can't open file 'MoviesItemOps02.py': [Errno 2] No such file or directory
ubuntu@ip-172-31-21-227:~$ python3 MoviesItemOps02.py
Get movie succeeded:
{'info': {'plot': 'Nothing happens at all.', 'rating': Decimal('0')},
 'title': 'The Big New Movie',
 'year': Decimal('2015'))
ubuntu@ip-172-31-21-227:~$ python3 MoviesItemOps03.py
Update movie succeeded:
{'Attributes': {'info': {'actors': ['Larry', 'Moe', 'Curly'],
                         'plot': 'Everything happens all at once.',
                         'rating': Decimal('5.5'))},
  'ResponseMetadata': {'HTTPHeaders': {'connection': 'keep-alive',
                                      'content-length': '156',
                                      'content-type': 'application/x-amz-json-1.0',
                                      'date': 'Tue, 13 Feb 2024 01:00:50 GMT',
                                      'server': 'Server',
                                      'x-amz-crc32': '3767510606',
                                      'x-amzn-requestid': 'AN60TUNANIUHLCB1AUJ0LG35JRVV4KQNS05AEMVJF66Q9ASUAAJG'},
                      'HTTPStatusCode': 200,
                      'RequestId': 'AN60TUNANIUHLCB1AUJ0LG35JRVV4KQNS05AEMVJF66Q9ASUAAJG',
                      'RetryAttempts': 0}}
ubuntu@ip-172-31-21-227:~$ 

```

*response received from the update operation*

```

dynamodb = boto3.resource('dynamodb', region_name='us-east-1')
table = dynamodb.Table('Movies')

try:
    response = table.delete_item(
        Key={
            'year': year,           } specifies the primary key attributes to identify the item to be deleted
            'title': title
        },
        ConditionExpression="info.rating <= :val",
        ExpressionAttributeValues={
            ":val": Decimal(rating)
        }
    )
except ClientError as e:
    if e.response['Error']['Code'] == "ConditionalCheckFailedException":
        print(e.response['Error']['Message'])
    else:
        raise
else:
    return response

if __name__ == '__main__':
    print("Attempting a conditional delete...")
    delete_response = delete_underrated_movie("The Big New Movie", 2015, 10)
    if delete_response:
        print("Delete movie succeeded:")
        pprint(delete_response)

```

Run the script:

```
$ python3 MoviesItemOps04.py
```

## 2.5 Query the data

To query data from a table, we use the `query` method. We must supply the `partition key` to the query method. The `sort key` is optional. For example, only supplying the year will return all movies of that year. Supplying both the year and the title will retrieve that specific movie. Further conditions can be applied to create various kinds of queries.

Let's look at a couple of examples:

### Example 1 Get all movies released in a year

Read the following code which accomplishes this task. This is available in the script `MoviesQuery01.py`.

```

import boto3
from boto3.dynamodb.conditions import Key
a class that helps construct conditions for DynamoDB queries

def query_movies(year, dynamodb=None):
    if not dynamodb:
        dynamodb = boto3.resource('dynamodb', region_name="us-east-1")

    table = dynamodb.Table('Movies')
    response = table.query(
        KeyConditionExpression=Key('year').eq(year)
    ) looking for items where the 'year' attribute is equal to the provided year when calling the function
    return response['Items'] return the list of items retrieved from the query

```

```

ubuntu@ip-172-31-21-227:~ $ python3 MoviesItemOps04.py
Attempting a conditional delete...
Delete movie succeeded: as rating of the movie is less than 10
HTTPStatus: 200, RequestId: 'AN60TUNANIUHLCB1AUJ0LG35JRVV4KQNS05AEMVJF66Q9ASUAAJG', RetryAttempts: 0
ubuntu@ip-172-31-21-227:~ $ python3 MoviesItemOps04.py
Attempting a conditional delete...
Delete movie succeeded: as rating of the movie is less than 10
HTTPStatus: 200, RequestId: 'JLKJJG18RCADTIDC4VJKQSNNNVV4KQNS05AEMVJF66Q9ASUAAJG', RetryAttempts: 0
ubuntu@ip-172-31-21-227:~ $

```

```

ubuntu@ip-172-31-21-227:~ $ python3 MoviesQuery01.py
Movies from 1985
1985 : A Nightmare on Elm Street Part 2: Freddy's Revenge
1985 : A View to a Kill
1985 : After Hours
1985 : Back to the Future
1985 : Better Off Dead...
1985 : Cocoon
1985 : Clue
1985 : Cocoon
1985 : Commando
1985 : Day of the Dead
1985 : Empire Mine
1985 : European Vacation
1985 : Explorers
1985 : Flesh/Blood
1985 : Friday the 13th: A New Beginning
1985 : Friday the 13th: Part 4
1985 : Girls Just Want to Have Fun
1985 : Just One of the Guys
1985 : Ladyhawke
1985 : Lethal
1985 : Lifeforce
1985 : Mad Max Beyond Thunderdome
1985 : Mask
1985 : Pale Rider
1985 : The Hunt for Red October

```

```

table = dynamodb.Table('Movies')
response = table.query(
    KeyConditionExpression=Key('year').eq(year)
)
looking for items where the 'year' attribute is equal to the provided year when calling the function
return response['Items']return the list of items retrieved from the query
items are extracted from the 'Item' key in the 'response'

if __name__ == '__main__':
    query_year = 1985
    print(f"Movies from {query_year}")
    movies = query_movies(query_year)
    for movie in movies:
        print(movie['year'], ":", movie['title'])

```

Notice the use of the parameter KeyConditionExpression instead of ConditionExpression. Also notice the use of function Key. The Boto 3 SDK automatically constructs a ConditionExpression for you when you use the Key (or Attr) function imported from boto3.dynamodb.conditions. The Attr function will be useful when applying conditions on attributes rather than keys. You can also specify a ConditionExpression directly as a string instead, as done previously.

Run the script:

```
$ python3 MoviesQuery01.py
```

#### Example 2 Get all movies released in a year with certain titles

Read the following code which accomplishes this task. This is available in the script MoviesQuery02.py.

```

from pprint import pprint
import boto3
from boto3.dynamodb.conditions import Key

def query_and_project_movies(year, title_range, dynamodb=None):
    if not dynamodb:
        dynamodb = boto3.resource('dynamodb', region_name="us-east-1")

    table = dynamodb.Table('Movies')
    print("Get year, title, genres, and lead actor")

    # Expression attribute names can only reference items in the projection expression.
    response = table.query(
        ProjectionExpression="#yr, title, info.genres, info.actors[0]",removed keyword first element of list
        ExpressionAttributeNames={"#yr": "year"},specifies the attributes that should be included in the result - limit the attributes returned in the response to only those we are interested in
        KeyConditionExpression=
            Key('year').eq(year) & Key('title').between(title_range[0], title_range[1])
    )
    return response['Items']

if __name__ == '__main__':
    query_year = 1992
    query_range = ('A', 'L')
    print(f"Get movies from {query_year} with titles from "
          f"\"{query_range[0]} to {query_range[1]}\"")
    movies = query_and_project_movies(query_year, query_range)
    for movie in movies:
        print(f"\n{movie['year']} : {movie['title']}")

if __name__ == '__main__':
    query_year = 1992
    query_range = ('A', 'L')
    print(f"Get movies from {query_year} with titles from "
          f"\"{query_range[0]} to {query_range[1]}\"")
    movies = query_and_project_movies(query_year, query_range)
    for movie in movies:
        print(f"\n{movie['year']} : {movie['title']}")

```

```
pprint(movie['info'])
```

The query prints movies from 1992, whose titles begin with letters A – L.

Run the script:

```
$ python3 MoviesQuery02.py
```

```

1985 : Girls Just Want to Have Fun
1985 : Last Action of the Guys
1985 : Ladyhawke
1985 : Legend
1985 : Lifeforce
1985 : Mad Max Beyond Thunderdome
1985 : Pale Rider
1985 : Pee-wee's Big Adventure
1985 : Perfect
1985 : Police Academy 2: Their First Assignment
1985 : Predator: Return to Blood Part II
1985 : Re-Animator
1985 : Real Genius
1985 : Return to Oz
1985 : Rocky IV
1985 : Robin Hood: Prince of Thieves
1985 : Silverado
1985 : St. Elmo's Fire
1985 : Teen Wolf
1985 : The Black Cauldron
1985 : The Breakfast Club

```

```

ubuntu@ip-172-31-21-227: ~
1985 : Witness
ubuntu@ip-172-31-21-227: ~$ python3 MoviesQuery02.py
Get movies from 1992 with titles from A to L
Get year, title, genres, and lead actor

1992 : A Few Good Men
['actors': ['Tom Cruise'], 'genres': ['Crime', 'Drama', 'Mystery', 'Thriller']}

1992 : A League of Their Own
['actors': ['Tom Hanks'], 'genres': ['Comedy', 'Drama', 'Sport']}

1992 : A River Runs Through It
['actors': ['Craig Sheffer'], 'genres': ['Drama']}

1992 : Aladdin
['actors': ['Scott Weinger'], 'genres': ['Animation', 'Adventure', 'Comedy', 'Family', 'Fantasy', 'Musical', 'Romance']}

1992 : Alien 3
['actors': ['Sigourney Weaver'], 'genres': ['Action', 'Sci-Fi', 'Thriller']}

1992 : Army of Darkness
['actors': ['Bruce Campbell'], 'genres': ['Comedy', 'Fantasy', 'Horror']}

1992 : Batman Returns
['actors': ['Michael Keaton'], 'genres': ['Action', 'Fantasy']}

1992 : Beethoven
['actors': ['Charles Grodin'], 'genres': ['Comedy', 'Drama', 'Family']}

1992 : Bitter Moon
['actors': ['Hugh Grant'], 'genres': ['Drama', 'Romance', 'Thriller']}

1992 : Boomerang
['actors': ['Eddie Murphy'], 'genres': ['Comedy', 'Drama', 'Romance']}

1992 : Braintdead
['actors': ['Timothy Balme'], 'genres': ['Comedy', 'Horror']}

ubuntu@ip-172-31-21-227: ~

```

```

1992 : Como agua para chocolate
['actors': ['Marco Leonardi'], 'genres': ['Drama', 'Romance']}

1992 : Damage
['actors': ['Jeremy Irons'], 'genres': ['Drama', 'Romance']}

1992 : Death Becomes Her
['actors': ['Meryl Streep'], 'genres': ['Comedy', 'Fantasy']}

1992 : El mariachi
['actors': ['Carlos Gallardo'], 'genres': ['Action', 'Crime', 'Drama', 'Thriller']}

1992 : Encino Man
['actors': ['Sean Astin'], 'genres': ['Comedy', 'Fantasy']}

1992 : Far and Away
['actors': ['Tom Cruise'], 'genres': ['Adventure', 'Drama', 'Romance', 'Western']}

1992 : FernGully: The Last Rainforest
['actors': ['Samantha Mathis'], 'genres': ['Animation', 'Adventure', 'Family', 'Fantasy', 'Thriller']}

1992 : Forever Young
['actors': ['Mel Gibson'], 'genres': ['Drama', 'Romance', 'Sci-Fi']}

1992 : Glengarry Glen Ross
['actors': ['Al Pacino'], 'genres': ['Drama']}

1992 : Home Alone 2: Lost in New York
['actors': ['Macaulay Culkin'], 'genres': ['Adventure', 'Comedy', 'Crime', 'Family']}

1992 : Howards End
['actors': ['Anthony Hopkins'], 'genres': ['Drama', 'Romance']}

1992 : Jennifer Eight
['actors': ['Andy Garcia'], 'genres': ['Crime', 'Drama', 'Mystery', 'Thriller']]
```

The query prints movies from 1992, whose titles begin with letters A – L.

Run the script:

```
$ python3 MoviesQuery02.py
```

Read the output carefully and compare it with the string given to the ProjectionExpression parameter.

Now open MoviesQuery02.py and remove the first two parameters:

```
ProjectionExpression="#yr, title, info.genres, info.actors[0]",  
ExpressionAttributeNames={"#yr": "year"},
```

Re-run the script and inspect the output again. What has changed?

## 2.6 Scan the data

The query method has the limitation that it must be supplied with a specific partition key value. However, what if we wish to write a query that needs to fetch data from multiple partitions? For instance: Print all movies between (and including) the years 1950 and 1959.

For such queries, we use the scan method. It reads all the items of a table, from all the partitions, then applies a filter\_expression, which you provide, and returns only the elements matching your criteria.

The following data implements the query, ‘Print all movies between (and including) the years 1950 and 1959.’ The code is available in MoviesScan.py.

```
from pprint import pprint  
import boto3  
from boto3.dynamodb.conditions import Key  
  
def scan_movies(year_range, display_movies, dynamodb=None):  
    if not dynamodb:  
        dynamodb = boto3.resource('dynamodb', region_name='us-east-1')  
  
    table = dynamodb.Table('Movies')  
  
    #scan and get the first page of results  
    response = table.scan(FilterExpression=Key('year').between(year_range[0],  
year_range[1]),  
                           initial scan operation to retrieve movies with release years within the specified range  
                           data = response['Items']  
                           parameter passed to the 'scan_movies' function  
                           display_movies(data)  
    #continue while there are more pages of results  
    while 'LastEvaluatedKey' in response:  
        response = table.scan(FilterExpression=Key('year').between(year_range[0],  
year_range[1]), ExclusiveStartKey=response['LastEvaluatedKey'])  
        data.extend(response['Items'])  
        extend the 'data' list with the items from the current response  
        start the next scan from where the previous one left off
```

removed  
ProjectionExpression  
and  
ExpressionAttributeNames

```
'Actors': ['Timothy Dalton'],  
'genres': ['Adventure', 'Comedy', 'Crime', 'Family'])  
1992 : Howards End  
{'actors': ['Anthony Hopkins'], 'genres': ['Drama', 'Romance'])  
1992 : Jennifer Eight  
{'actors': ['Andy Garcia'], 'genres': ['Crime', 'Drama', 'Mystery', 'Thriller'])  
1992 : Juice  
{'actors': ['Omar Epps'], 'genres': ['Crime', 'Drama', 'Thriller'])  
ubuntu@ip-172-31-21-227:~$  
  
[{"genres": ["Adventure", "Comedy", "Crime", "Family"],  
"image_url": "http://ia.media-imdb.com/images/M/MV5BMTYyODQzNDRkLV5BM15BanBnXkFtZTcwOTM5NTU1MQ@@._V1_SX400_.jpg",  
"plot": "One year after Kevin was left home alone and had to defeat a pair of bumbling burglars, he accidentally finds himself in New York City, and the same criminals are not far behind.",  
"rank": Decimal("1279"),  
"rating": Decimal("7.4"),  
"release_date": "1992-11-15T00:00:00Z",  
"running_time_secs": Decimal("7200")},  
  
1992 : Howards End  
{"actors": ["Anthony Hopkins", "Emma Thompson", "Vanessa Redgrave"],  
"directors": ["James Ivory"],  
"genres": ["Drama", "Romance"],  
"image_url": "http://ia.media-imdb.com/images/M/MV5BOTE2NzI4MTU4NV5BM15BanBnXkFtZTcwMzY5ODIxMw@@._V1_SX400_.jpg",  
"plot": "A businessman thwarts his wife's bequest of a estate to another",  
"rank": Decimal("3513"),  
"rating": Decimal("7.4"),  
"release_date": "1992-03-13T00:00:00Z",  
"running_time_secs": Decimal("8400")},  
  
1992 : Jennifer Eight  
{"actors": ["Andy Garcia", "Uma Thurman", "Lance Henriksen"],  
"directors": ["Bruce Robinson"],  
"genres": ["Crime", "Drama", "Mystery", "Thriller"],  
"image_url": "http://ia.media-imdb.com/images/M/MV5BMTgxMTk3NTc1OV5BM15BanBnXkFtZTcwMzI5NzMyM2@@._V1_SX400_.jpg",  
"plot": "A man who has been away from his wife for 15 years returns home to find she has died and her new husband is his best friend.",  
"rank": Decimal("4338"),  
"rating": Decimal("6.1"),  
"release_date": "1992-11-06T00:00:00Z",  
"running_time_secs": Decimal("7440")},  
  
1992 : Juice  
{"actors": ["Omar Epps", "Tupac Shakur", "Jermaine 'Huggy' Hopkins"],  
"directors": ["Ernest R. Dickerson"],  
"genres": ["Crime", "Drama", "Thriller"],  
"image_url": "http://ia.media-imdb.com/images/M/MV5BMTI2MTg1OTgiOF5BM15BanBnXkFtZTcwMzI0NjMyM2@@._V1_SX400_.jpg",  
"plot": "It's about 4 innercity teens who get caught up in the pursuit of power and happiness, which they refer to as 'the juice'.",  
"rank": Decimal("4806"),  
"rating": Decimal("6.7"),  
"release_date": "1992-01-17T00:00:00Z",  
"running_time_secs": Decimal("5700")}  
ubuntu@ip-172-31-21-227:~$
```

printing all info

```

    display_movies(data)

    return data

if __name__ == '__main__':
    def print_movies(movies):
        for movie in movies:
            #print(f"\n{movie['year']} : {movie['title']}")
            #pprint(movie['info'])
            pprint(movie)

    query_range = (1950, 1959)
    print(f"Scanning for movies released from {query_range[0]} to {query_range[1]}...")
    scan_movies(query_range, print_movies)

```

Run the script:

```
$ python3 MoviesScan.py
```

There are a few important things to understand in this code. First is the `FilterExpression` parameter. Here the `Key` function has been used in the same way it was used earlier to create `KeyConditionExpression`.

The other thing to note is how we retrieve the results of the `scan` function. This has to do with the concept of pagination in DynamoDB. DynamoDB applies and returns the results of queries and scans one page at a time where a page can contain at max 1MB of data. Therefore, the first call to `scan` returns the first page of results. The while loop then continues to re-call the `scan` function as long as there are more pages to filter results from. The response object maintains an attribute `LastEvaluatedKey` which is `None` if there are more pages following the current one. The scanning process always begins from the key pointed to by the `ExclusiveStartKey` parameter. The expression `ExclusiveStartKey=response['LastEvaluatedKey']` points the `ExclusiveStartKey` to the top of the next page, and so on.

There is a [paginator module in Boto3](#) which provides abstractions for doing pagination in different scenarios.

*For more information on queries and scans, visit the links below:*

[Working with Queries in DynamoDB](#)

[Working with Scans in DynamoDB](#)

## 2.7 Delete table

Read the following code which deletes a table from the database. This code is available in `MoviesDeleteTable.py`

```

ubuntu@ip-172-31-21-227: ~
{
    'rating': Decimal('7.7'),
    'release_date': '1956-04-12T00:00:00Z',
    'running_time_secs': Decimal('91800'),
    'title': 'The Man in the Gray Flannel Suit',
    'year': Decimal('1956')
}
['info': {'actors': ['John Wayne', 'Jeffrey Hunter', 'Vera Miles'],
          'directors': ['John Ford'],
          'plot': "A Civil War veteran embarks on a journey to rescue his niece from an Indian tribe",
          'rating': Decimal('3090'),
          'release_date': '1956-03-13T00:00:00Z',
          'running_time_secs': Decimal('7140')},
 'title': 'The Searchers',
 'year': Decimal('1956')
}
['info': {'actors': ['Montgomery Clift', 'Charlton Heston', 'Yul Brynner', 'Anne Baxter'],
          'directors': ['Cecil B. DeMille'],
          'plot': "An Egyptian Prince, Moses, learns of his true heritage as a Hebrew and his divine mission as the deliverer of his people",
          'rating': Decimal('2340'),
          'release_date': '1956-01-10-05T00:00:00Z',
          'running_time_secs': Decimal('13200')},
 'title': 'The Ten Commandments',
 'year': Decimal('1956')
}
['info': {'actors': ['Hugh Marlowe', 'Nancy Gates', 'Nelson Leigh'],
          'directors': ['Richard Brooks'],
          'plot': "Astronauts returning from a voyage to mars are caught in a time warp and are propelled into a post-Apocalyptic Earth",
          'rating': Decimal('1992'),
          'release_date': '1956-03-25T00:00:00Z',
          'running_time_secs': Decimal('4800')},
 'title': 'World Without End',
 'year': Decimal('1956')
}
ubuntu@ip-172-31-21-227: ~

```

```
import boto3

def delete_movie_table(dynamodb=None):
    if not dynamodb:
        dynamodb = boto3.resource('dynamodb', region_name="us-east-1")

    table = dynamodb.Table('Movies')
    table.delete()

    if __name__ == '__main__':
        delete_movie_table()
        print("Movies table deleted.")
```

Run the script:

```
$ python3 MoviesDeleteTable.py
```

## 2.8 (Optional) Practice Exercises

Recreate the table `Movies` (if you have deleted it), and load the data into it.

```
$ python3 MoviesCreateTable.py
$ python3 MoviesLoadData.py
```

Write python scripts to perform the following operations on the movies database:

1. Print the titles of all movies released in 1994.
2. Print complete information on the movie 'After Hours' released in 1985.
3. Print all movies released before 2000.
4. Print only the years and titles of movies starring Tom Hanks.
5. Remove all movies released before 2000.

### Sources

1. Amazon Dynamo DB Developer Guide for Python

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GettingStarted.Python.html>

```
'year': Decimal('1956')}

ubuntu@ip-172-31-21-227:~$ python3 MoviesDeleteTable.py
Movies table deleted.

ubuntu@ip-172-31-21-227:~$
```